

RediSearch



Agenda

- Why??
- Example Use Case
- Autocompletion
- Loading Data
- Searching
- Aggregating
- RedisLabs Extras

Search Database vs RDBMS

Search

RDBMS

Where's my data?

Give me *this* data

Index data but also store

Store data but also index

Built for de-normalized data

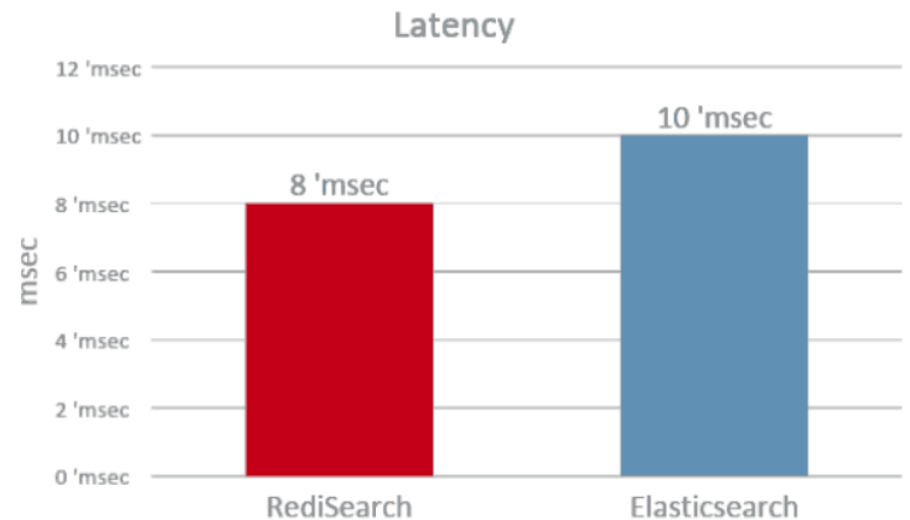
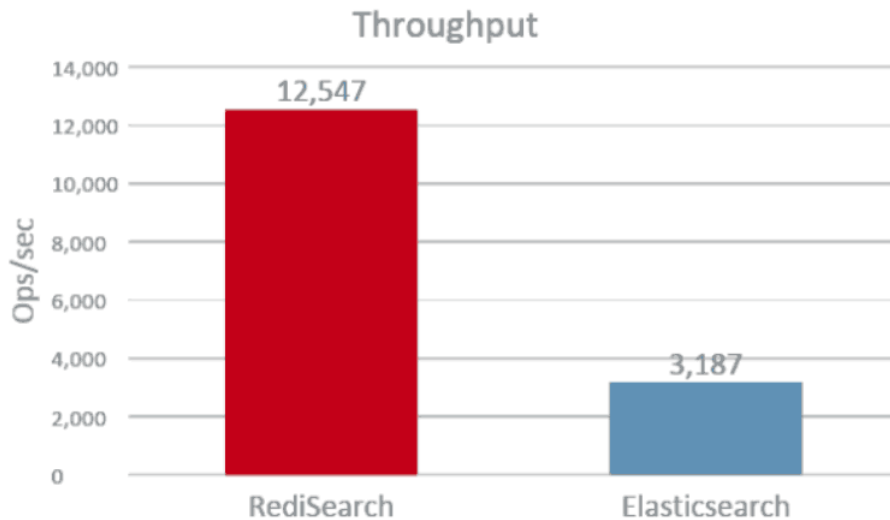
Optimized for normalized data

AP in CAP Theorum

CP in CAP Theorum

Why Search in Redis??

Speed !!



Why Search in Redis??

I want my data Now !!

Documents are available for reading *immediately*

Why Search in Redis??

Easy !!

Load the RediSearch module and go No barriers to
searching *now*

Example Use Case

Fortune 500 Companies

- Rank
- CEO
- Sector
- Industry
- Ticker

Set Up Autocompletion

```
> FT.SUGLEN ac
0
> FT.SUGADD ac Walmart 1.0
1
> FT.SUGADD ac "Berkshire Hathaway" 1.0
2
> FT.SUGADD ac "Apple" 1.0
3
...
> FT.SUGLEN ac
500
```


Query the Autocompleter

```
> FT.SUGGET ac birk  
(empty list or set)  
  
# No matches - Let's try Fuzzy matching  
> FT.SUGGET ac birk FUZZY  
1) "Berkshire Hathaway"
```

Code Example

```
from redisearch import AutoCompleter, Suggestion
ac = AutoCompleter(
    'ac',
    conn = client.redis
)

if ac.len() < 1:
    load_data()

for row in csv_reader:
    ac.add_suggestions(Suggestion(row[1], 1.0))
```

Loading Data

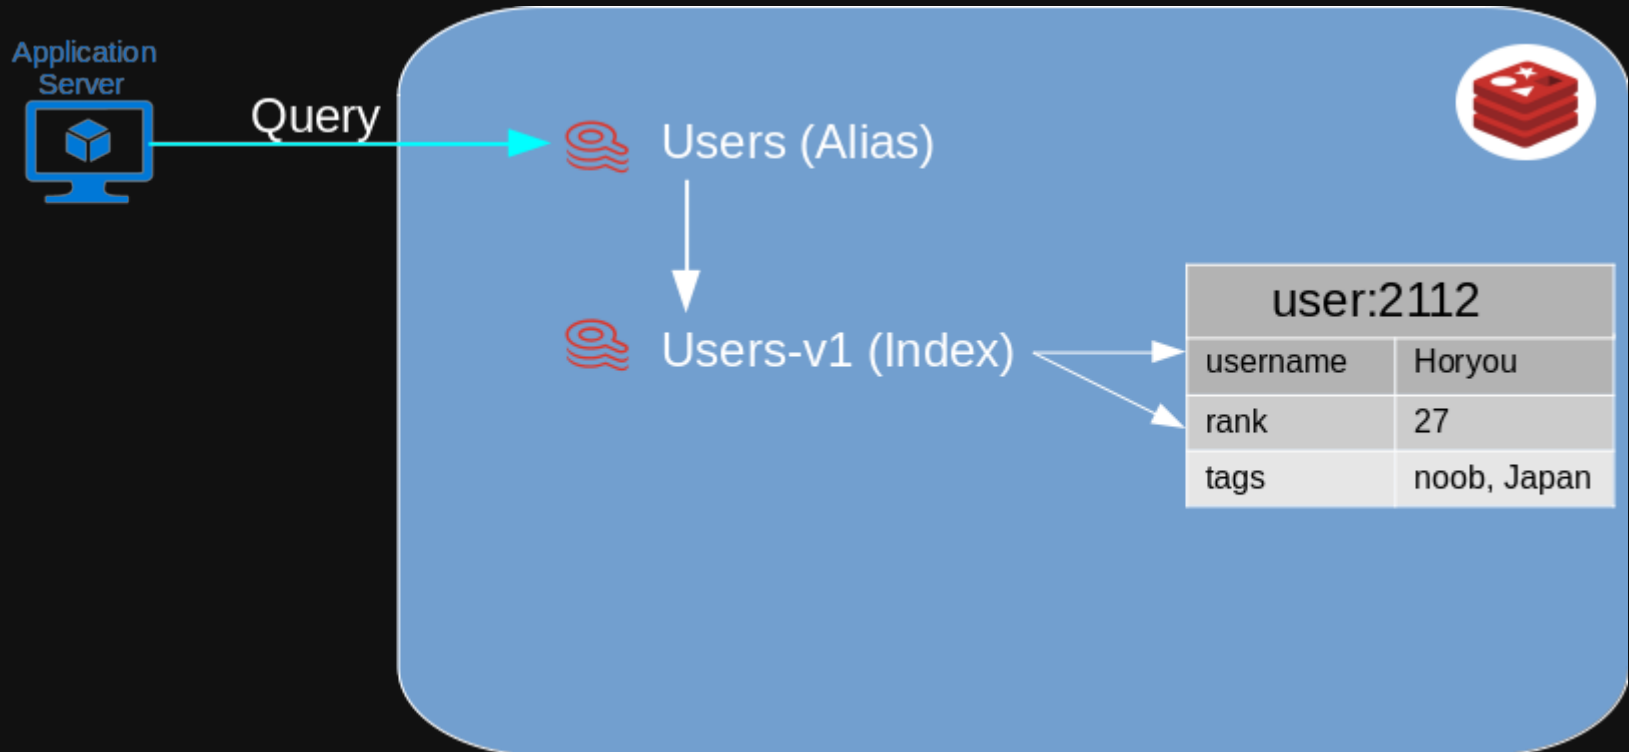
Pro-Tip

Consider aliasing your Index
before you need to re-index your data

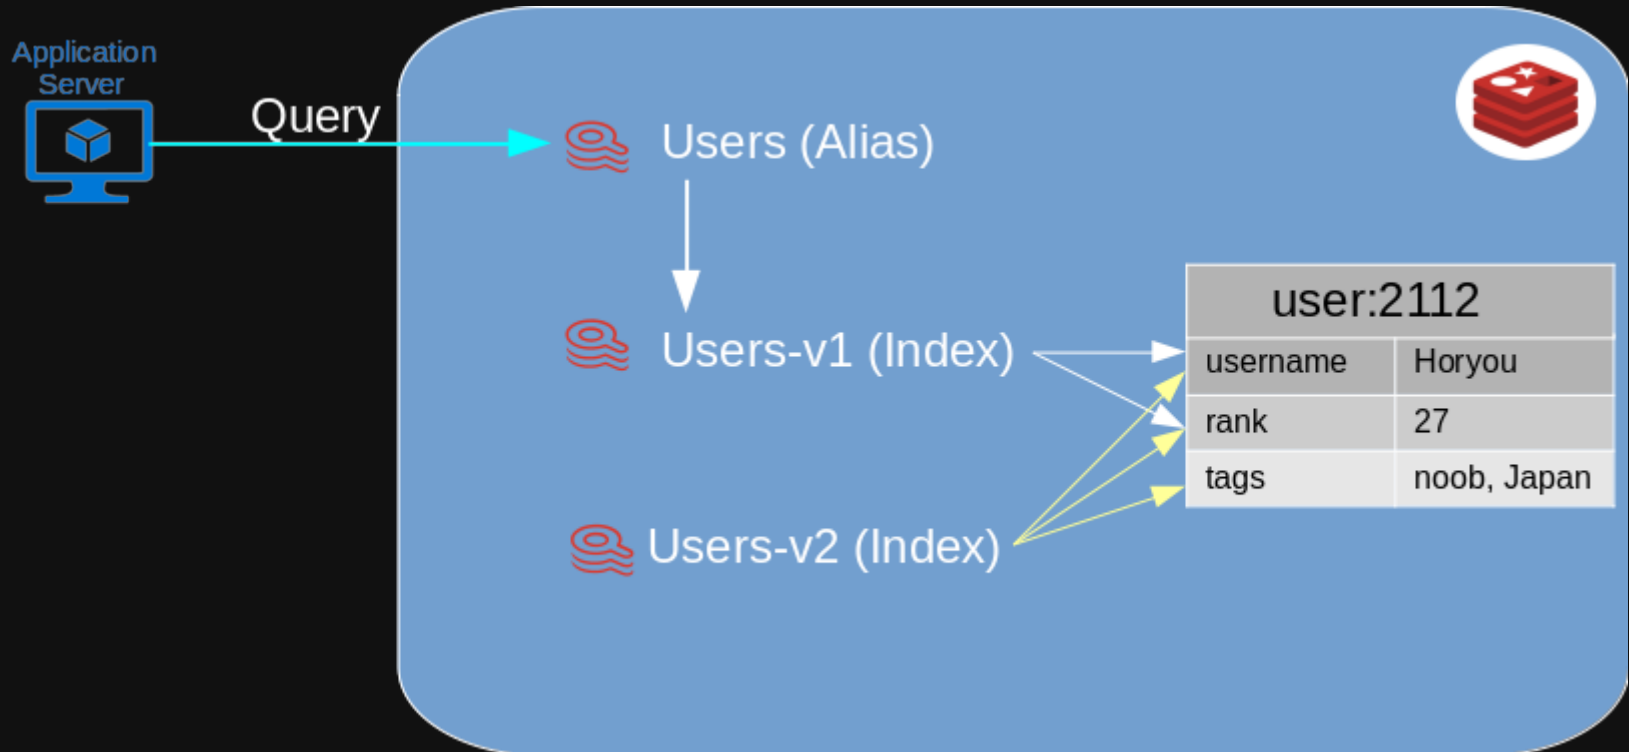
Why Alias?

- I want to modify my schema without application changes
- I want to future proof my application

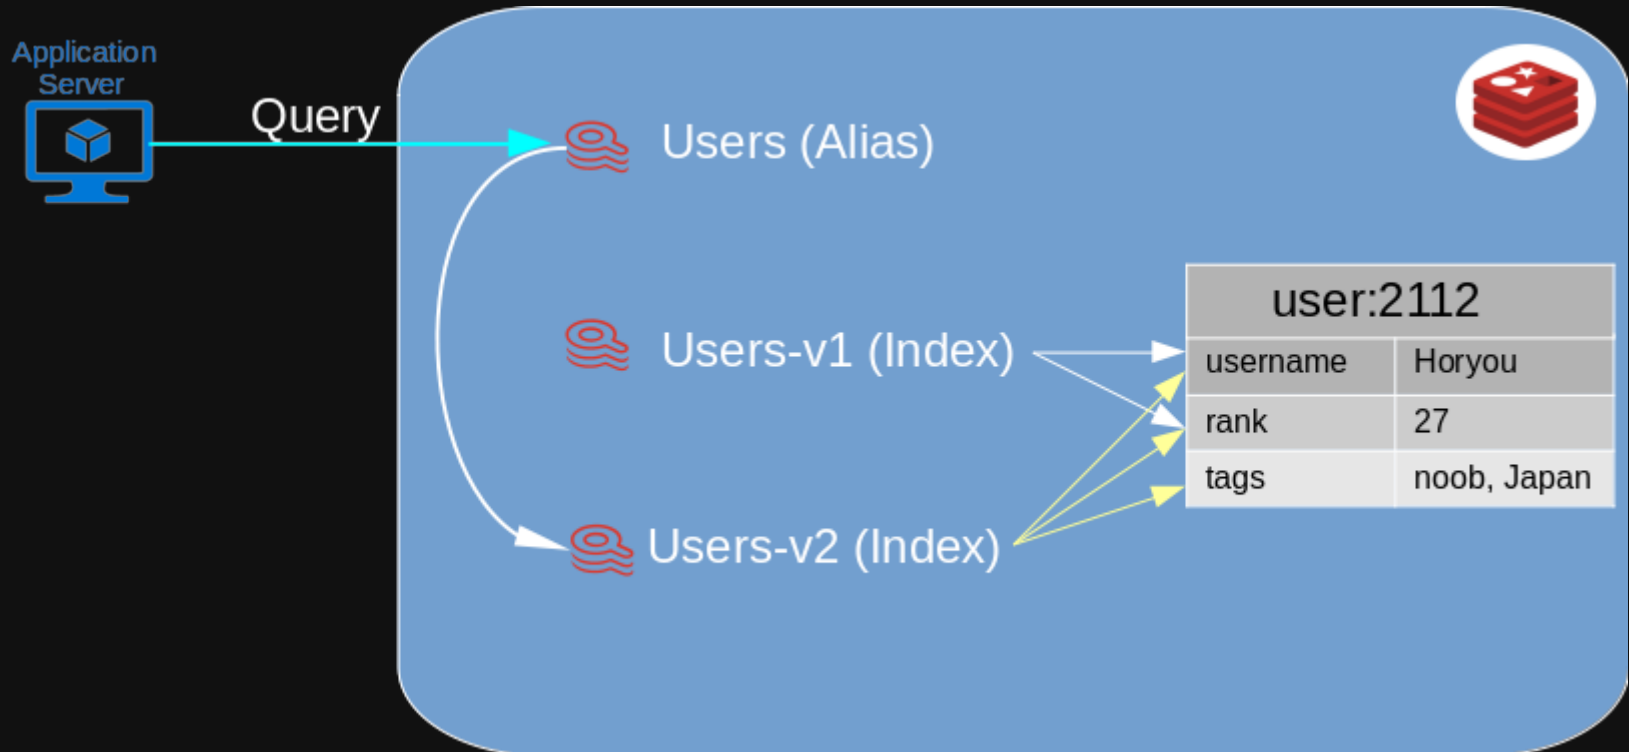
Alias: Initial




Alias: Create New Index with tags too



Alias: Flip Alias to New Index



Alias: Remove old index



By default, FT.DROPINDEX **does not delete** the document hashes associated with the index.

Adding the DD option deletes the hashes as well.

Create our Schema

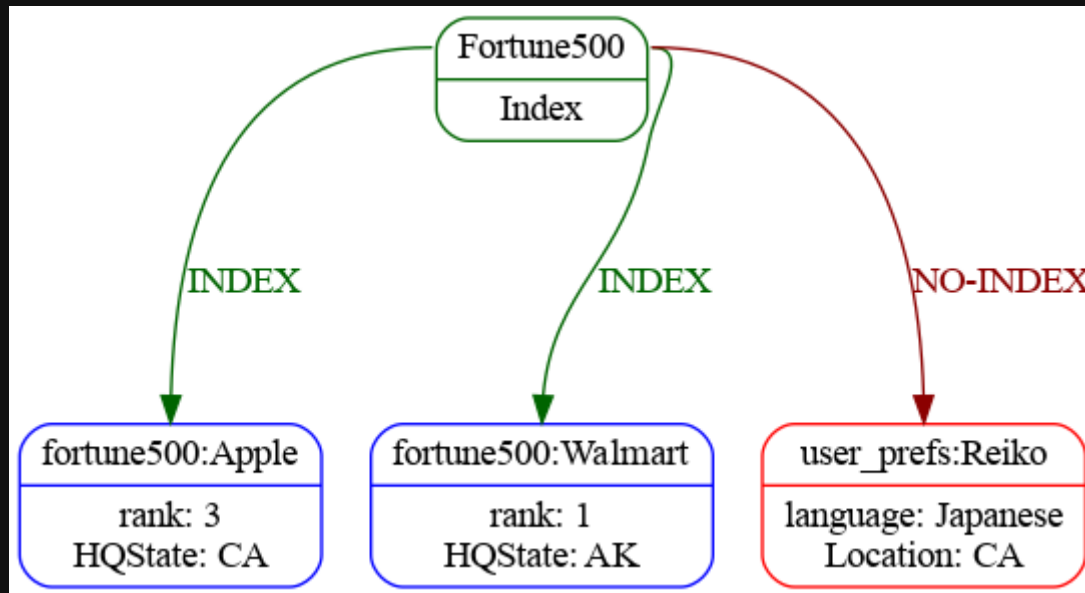
```
FT.CREATE fortune500
  ON HASH PREFIX 1 fortune500:
  LANGUAGE English
  SCORE_FIELD title SCORE 0.5
  SCHEMA
    title TEXT WEIGHT 5.0
    employees NUMERIC SORTABLE
```

Schema breakdown

```
FT.CREATE INDEX_NAME  
  ON HASH PREFIX <COUNT> PREFIX1 PREFIX2
```

```
FT.CREATE myIndex  
  ON HASH PREFIX 3 end_user: admin_user: ro_user:
```

Prefixes?!?



- Writing applications do not have to be search aware
- Add search indices on “legacy” data
- Flexibility to add new fields or un-indexed fields

Languages

```
FT.CREATE fortune500
...
LANGUAGE English
...
```

Arabic	Danish	Dutch	English	Finnish
French	German	Hungarian	Italian	Norwegian
Portuguese	Romanian	Russian	Spanish	
Swedish	Tamil	Turkish	Chinese	

Languages - Why?

Stemming

- search going will return “going”, “go”, “gone”
- plurals, verb forms
- suffixes eg: ational ation ator: replace by ate
- custom dictionaries are possible

Scoring

```
FT.CREATE fortune500
...
SCORE_FIELD title SCORE 1.0
SCORE_FIELD industry SCORE 0.5
...
```

Scoring

```
> FT.SEARCH fortune500 "%computer%" WITHSCORES
2) "fortune500:ComputerSciences"
3) "2.5"
4) 1) "rank"
   2) "379"
   3) "website"
   4) "http://www.dxc.technology"
   ...
5) "fortune500:ActivisionBlizzard"
6) "0.5"
7) 1) "rank"
   2) "406"
   ...
10) "Computer Software"
```


Searching

https://oss.redislabs.com/redisearch/Query_Syntax/

Searching

Range Match

```
@assets:[20000, 30000]
```

Negative Match

```
-@hqstate:(NY|CA)
```

Searching

Optional Match

```
~@sector:\"Health Care\"
```

Fuzzy Match

```
~@ceo:%%Brian%%
```

Searching - Find Everything

```
> FT.SEARCH fortune500 "*" LIMIT 0 1
1) (integer) 500
2) "fortune500:TollBrothers"
3) 1) "rank"
    2) "497"
    3) "website"
    4) "http://www.tollbrothers.com"
    5) "employees"
    6) "4200"
    7) "sector"
    8) "Engineering & Construction"
    9) "industry"
    ...
```

Searching - Query by field

```
> FT.SEARCH fortune500 "@assets:[20000, 30000]
    -@hqstate:NY ~@sector:\"Health Care\"" WITHSCORES
2) "fortune500:CommunityHealthSystems"
3) "0.2000000000000000001"
4) 1) "rank"
    2) "130"
    7) "sector"
    8) "Health Care"
29) "fortune500:Level3Communications"
30) "0.1000000000000000001"
31) 1) "rank"
    2) "336"
    7) "sector"
    8) "Telecommunications"
```

Searching - Code

```
from redisearch import Client, Query
client = Client('fortune500')
client.search(Query("technology").limit_fields('sector')).docs
```

Searching - Tags

Pros

Easy to add tags

Super efficient search

Low storage requirements

Cons

Low cardinality is critical

Only exact matches

Aggregations

<https://oss.redislabs.com/redisearch/Aggregations/>

Aggregations

```
> FT.AGGREGATE fortune500 "*" GROUPBY 1 @hqstate
  REDUCE COUNT 0 AS my_count SORTBY 2 @my_count DESC LIMIT 0 2
1) (integer) 37
2) 1) "hqstate"
   2) "ny"
   3) "my_count"
   4) "54"
3) 1) "hqstate"
   2) "ca"
   3) "my_count"
   4) "53"
```

Aggregations with functions

```
> FT.AGGREGATE fortune500 "*" APPLY "upper(@hqstate)" AS state
GROUPBY 1 @state REDUCE COUNT 0 AS my_count
SORTBY 2 @my_count DESC LIMIT 0 2
1) (integer) 37
2) 1) "state"
   2) "NY"
   3) "my_count"
   4) "54"
3) 1) "state"
   2) "CA"
   3) "my_count"
   4) "53"
```

Aggregations with functions

```
from redisearch import Client, aggregation, reducers
client = Client('fortune500')
ar = aggregation.AggregateRequest().
    group_by("@hqcity", reducers.count().alias('my_count')).
    sort_by(aggregation.Desc('@my_count'))
client.aggregate(ar).rows
```

Aggregations Caveat



RedisLabs Extras



Redis Insight

<https://redislabs.com/redisinsight/>

Free!

redisinsight

search ▾

OVERVIEW

BROWSE

Browser

CLI

Streams

RedisGraph

RedisGears β

RedisTimeSeries

RediSearch

ANALYSE

Memory Analysis ▾

Profiler

Slowlog

BULK ACTIONS β

DATABASE

Configuration

Client List

SETTINGS

AGGREGATE ▾

fortune500 ▾

Execute Query ▶

Use Shift+Enter to format your query in multiple lines

FT.AGGREGATE fortune500: "*" GROUPBY 2 @hqstate @industry REDUCE COUNT 0 AS my_count SORTBY 4 @... [] ↺ ^ (x)

my_count

6

tx

petroleum refining

6

ny

commercial banks

6

ny

diversified financials

6

ny

insurance: property and casualty (stock)

6

ca

semiconductors and other electronic components

6

tx

pipelines

5

tx

mining, crude-oil production

5

mi

motor vehicles and parts

5

ca

computer software

5

ca

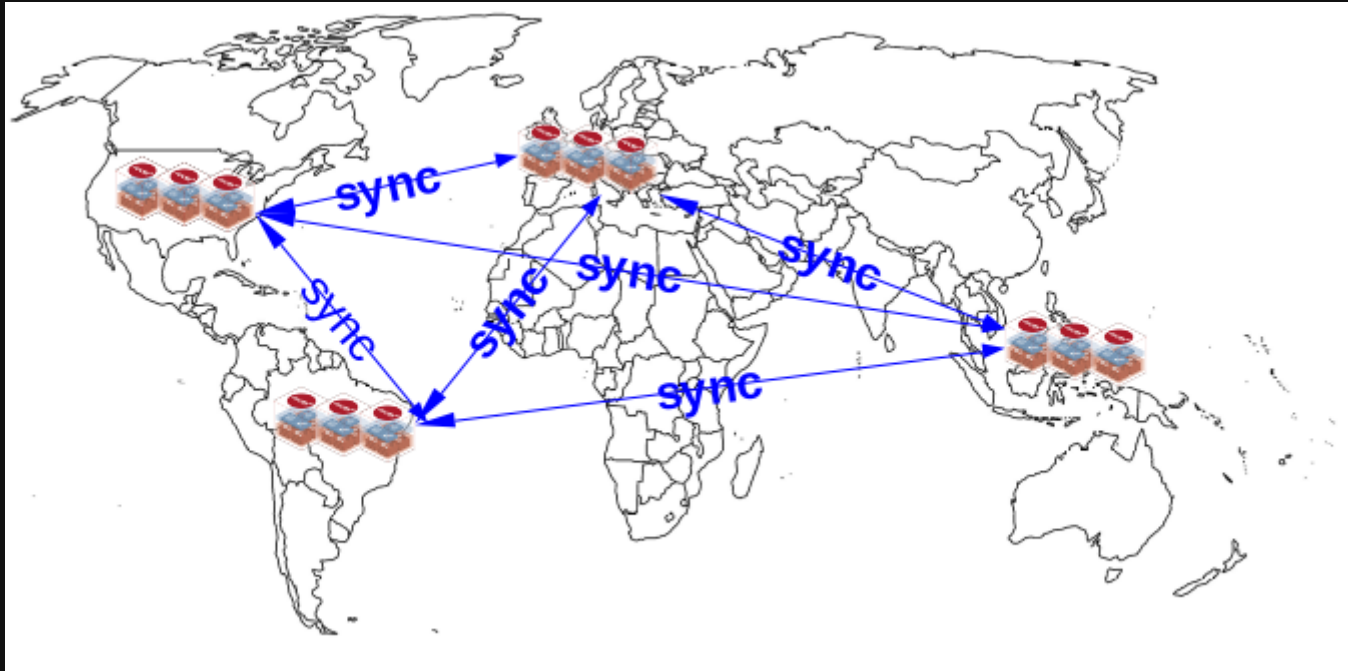
internet services and retailing

5

Total Documents Matched: 365

See More

Redis Enterprise Active/Active



- Based off of CRDTs
- Used for
 - Disaster Recovery
 - Data Migration
 - Enhanced Uptime

Thank you